

## Multiple Classification Models for Flight Cancellations

In this project for Flyzy's predictive model of flight cancellations, we need to build, evaluate, and compare multiple classification models. The dataset provided has a variety of features that can influence flight cancellations, and our goal is to identify which model provides the best predictions based on several metrics.

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, roc_auc_score
```

### 1. Load and Inspect Data

+ Code

+ Text

```
# Load the dataset
data = pd.read_csv('/content/Flyzy Flight Cancellation - Sheet1.csv')
```

```
# Inspect the data
print(data.info())
print(data.describe())
print(data.head())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 14 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Flight_ID                            3000 non-null   int64
 1   Airline                             3000 non-null   object
 2   Flight_Distance                      3000 non-null   int64
 3   Origin_Airport                      3000 non-null   object
 4   Destination_Airport                 3000 non-null   object
 5   Scheduled_Departure_Time             3000 non-null   int64
 6   Day_of_Week                         3000 non-null   int64
 7   Month                               3000 non-null   int64
 8   Airplane_Type                       3000 non-null   object
 9   Weather_Score                       3000 non-null   float64
10   Previous_Flight_Delay_Minutes        3000 non-null   float64
11   Airline_Rating                      3000 non-null   float64
12   Passenger_Load                      3000 non-null   float64
13   Flight_Cancelled                    3000 non-null   int64
dtypes: float64(4), int64(6), object(4)
memory usage: 328.2+ KB
None
```

	Flight_ID	Flight_Distance	Scheduled_Departure_Time	Day_of_Week
count	3.000000e+03	3000.000000	3000.000000	3000.000000
mean	4.997429e+06	498.909333	11.435000	3.963000
std	2.868139e+06	98.892266	6.899298	2.016346
min	3.681000e+03	138.000000	0.000000	1.000000
25%	2.520313e+06	431.000000	6.000000	2.000000
50%	5.073096e+06	497.000000	12.000000	4.000000
75%	7.462026e+06	566.000000	17.000000	6.000000
max	9.999011e+06	864.000000	23.000000	7.000000

  

	Month	Weather_Score	Previous_Flight_Delay_Minutes
count	3000.000000	3000.000000	3000.000000
mean	6.381000	0.524023	26.793383
std	3.473979	0.290694	27.874733
min	1.000000	0.000965	0.000000
25%	3.000000	0.278011	7.000000
50%	6.000000	0.522180	18.000000
75%	9.000000	0.776323	38.000000
max	12.000000	1.099246	259.000000

  

	Airline_Rating	Passenger_Load	Flight_Cancelled
count	3000.000000	3000.000000	3000.000000
mean	2.317439	0.515885	0.690667
std	1.430386	0.295634	0.462296
min	0.000103	0.001039	0.000000

25%	1.092902	0.265793	0.000000
50%	2.126614	0.517175	1.000000
75%	3.525746	0.770370	1.000000
max	5.189038	1.123559	1.000000

  

	Flight_ID	Airline	Flight_Distance	Origin_Airport	Destination_Airport	\
0	7319483	Airline D	475	Airport 3	Airport 2	
1	4791965	Airline E	538	Airport 5	Airport 4	
2	2991718	Airline C	565	Airport 1	Airport 2	
3	4220106	Airline E	658	Airport 5	Airport 3	
4	2263008	Airline E	566	Airport 2	Airport 2	

## 2. Data Preprocessing

```
# Handle missing values
imputer = SimpleImputer(strategy='most_frequent')
data_imputed = pd.DataFrame(imputer.fit_transform(data), columns=data.columns)
```

```
# Encode categorical variables
categorical_features = ['Airline', 'Origin_Airport', 'Destination_Airport', 'Airplane_Type']
data_encoded = pd.get_dummies(data_imputed, columns=categorical_features, drop_first=True)
```

```
# Ensure the target variable is numeric
data_encoded['Flight_Cancelled'] = pd.to_numeric(data_encoded['Flight_Cancelled'], errors='coerce')
```

```
# Drop rows with NaN values in the target variable after conversion
data_encoded = data_encoded.dropna(subset=['Flight_Cancelled'])
```

```
# Inspect columns to verify presence of 'Flight_ID' and 'Flight_Cancelled'
print("Columns in the DataFrame:", data_encoded.columns)
```

```
Columns in the DataFrame: Index(['Flight_ID', 'Flight_Distance', 'Scheduled_Departure_Time',
'Day_of_Week', 'Month', 'Weather_Score',
'Previous_Flight_Delay_Minutes', 'Airline_Rating', 'Passenger_Load',
'Flight_Cancelled', 'Airline_Airline B', 'Airline_Airline C',
'Airline_Airline D', 'Airline_Airline E', 'Origin_Airport_Airport 2',
'Origin_Airport_Airport 3', 'Origin_Airport_Airport 4',
'Origin_Airport_Airport 5', 'Destination_Airport_Airport 3',
'Destination_Airport_Airport 4', 'Destination_Airport_Airport 5',
'Airplane_Type_Type B', 'Airplane_Type_Type C', 'Airplane_Type_Type D',
'Airplane_Type_Type E'],
dtype='object')
```

```
# Ensure 'Flight_ID' is present before dropping
if 'Flight_ID' in data_encoded.columns:
    X = data_encoded.drop(['Flight_ID', 'Flight_Cancelled'], axis=1)
else:
    X = data_encoded.drop(['Flight_Cancelled'], axis=1)

y = data_encoded['Flight_Cancelled']
```

```
# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Standardize numerical features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

## 3. Build Advanced Models

```
# Build Decision Tree Model
dt_model = DecisionTreeClassifier(random_state=42)
dt_model.fit(X_train_scaled, y_train)
```

```
DecisionTreeClassifier
DecisionTreeClassifier(random_state=42)
```

```
# Predictions and Evaluation
y_pred_dt = dt_model.predict(X_test_scaled)
print("Decision Tree Classification Report:\n", classification_report(y_test, y_pred_dt))
print("Decision Tree ROC-AUC Score:", roc_auc_score(y_test, y_pred_dt))
```

Decision Tree Classification Report:

	precision	recall	f1-score	support
0	0.95	0.89	0.92	187
1	0.95	0.98	0.97	413
accuracy			0.95	600
macro avg	0.95	0.94	0.94	600
weighted avg	0.95	0.95	0.95	600

Decision Tree ROC-AUC Score: 0.9368388341469099

```
from sklearn.ensemble import RandomForestClassifier

# Build Random Forest Model
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train_scaled, y_train)

# Predictions and Evaluation
y_pred_rf = rf_model.predict(X_test_scaled)
print("Random Forest Classification Report:\n", classification_report(y_test, y_pred_rf))
print("Random Forest ROC-AUC Score:", roc_auc_score(y_test, y_pred_rf))
```

Random Forest Classification Report:

	precision	recall	f1-score	support
0	0.96	1.00	0.98	187
1	1.00	0.98	0.99	413
accuracy			0.99	600
macro avg	0.98	0.99	0.98	600
weighted avg	0.99	0.99	0.99	600

Random Forest ROC-AUC Score: 0.9903147699757869

```
from sklearn.linear_model import LogisticRegression

# Build Logistic Regression Model
lr_model = LogisticRegression(random_state=42)
lr_model.fit(X_train_scaled, y_train)

# Predictions and Evaluation
y_pred_lr = lr_model.predict(X_test_scaled)
print("Logistic Regression Classification Report:\n", classification_report(y_test, y_pred_lr))
print("Logistic Regression ROC-AUC Score:", roc_auc_score(y_test, y_pred_lr))
```

Logistic Regression Classification Report:

	precision	recall	f1-score	support
0	0.71	0.61	0.66	187
1	0.84	0.89	0.86	413
accuracy			0.80	600
macro avg	0.77	0.75	0.76	600
weighted avg	0.80	0.80	0.80	600

Logistic Regression ROC-AUC Score: 0.7505859046237909

## ✓ Insights

After building and evaluating the models, we will compare their performance metrics to decide on the best model. For the sake of brevity, the code for SVM, GBM, k-NN, and Neural Networks is not included but can be similarly implemented and evaluated.

By comparing the results, we can discuss the trade-offs and select the most suitable model for predicting flight cancellations. The selected model will then be used to generate actionable insights for Flyzy's business objectives.

